

# Lab 3: TCP Congestion and UDP Socket Programming

50.012 Networks

Hand-out: September 28  
eDimension hand-in: October 5

## 1 Objectives

- Experiment with buffer sizes, network scheduling, and TCP congestion window
- Develop a simple client-server application based on UDP sockets
- Goal: Understand why larger buffers can be disadvantageous
- Part of this exercise is based on the Stanford CS244 class lab1

## 2 Setup

### 2.1 Lab3 setup

- Download the lab3.zip from eDimension and unpack to some local folder, e.g. ~/lab3/
- Change into the lab3 folder, and open up another terminal tab with CTRL+SHIFT+T

### 2.2 mininet

- This exercise assumes that you have a running mininet installation
  - It should be installed on the lab machines already
  - On (x)ubuntu, you can install it by running our `install.sh`. Mininet can run on many \*nixes, but NOT on Windows or Mac

#### 1. What is Mininet?

- Mininet is a network emulator (mostly written in Python), that allows you to define network topologies and to connect emulated hosts.
- The emulated hosts run on your OS, and can execute your client/server applications easily.
- You can also easily change link parameters such a packet loss rate, delay, and bandwidth.
- Mininet is mainly intended to experiment with

*Software Defined Networks and OpenFlow*, but is also convenient for this lab session. We might look at OpenFlow later in the term.

- More details at <http://mininet.org>, in particular <http://mininet.org/walkthrough/>
- If you like videos: some basic things of this sheet are also demo'ed in <https://www.youtube.com/watch?v=jmlgXaocwiE>

## 2. Installation and first test

- Start up mininet:

```
sudo ./run.sh
```

- Get familiar with mininet commands using mininet's help system:

```
mininet> help [topic]
```

- Use the following mininet command to spawn two terminals, belonging to two nodes in the emulated network.

```
mininet> xterm h1 h2
```


- Using these terminals, you can run applications or perform ping operations on the emulated hosts
  - To close running commands like `ping -c 100`, try CTRL-C
- Find out the IP addresses of h1 and h2
- Note that `mininet` allows to run a command from its command prompt as if the command is run from any virtual host. For example, if we want to ping h2 from h1 we can run:

```
mininet> h1 ping h2
```

- To leave/close Mininet, try CTRL-D on the main command line


## 3 Queue Length and TCP

In this exercise we will study the dynamics of TCP in home networks. Take a look at the figure below which shows a "typical" home network with a "Home Router" connected to an end host. The "Home Router" is connected via Cable or DSL to a Headend router at the Internet access provider's office. We are going to study what happens when we download data from a remote server to the End Host in this home network.



Problem overview

In a real network it's hard to measure the TCP congestion window  $cwnd$  (because it's private to the Server) and the buffer occupancy (because it's private to the router). Luckily, we have mininet to help allow us to get these values easily.



### 3.1 Simple TCP connection (wget)

- Start up mininet

```
sudo ./run.sh
```

- In the running mininet emulation, test how long it takes to download a webpage on h1 from h2.
- wget's visualization is a bit confusing. You can find the total

time by subtracting start time from end time, or by looking at the time at end of last part line.

```
mininet> h2 wget h1
```

- We discussed the TCP congestion window in the last week, what is your expectation how the  $cwnd$  changes during this HTTP session?
- How big is the router queue/buffer in your opinion?

### 3.2 With parallel load (iperf)

- To see how the dynamics of a long flow differs from a short flow (which never leaves slow-start), we are going to repeat Part 2 for a “streaming video flow”. Instead of actually watching videos on your machine, we are going to set up a long-lived high speed TCP connection instead, to emulate a long-lived video flow.
- You can generate long flows using the `iperf` command. `iperf` is a tool for performing network throughput measurements. It can test either TCP or UDP throughput. To perform an `iperf` test the user must establish both a server (to discard traffic) and a client (to generate traffic). We have wrapped `iperf` in a script that you can run as follows:

```
mininet> h1 ./iperf.sh
```

- You can see the throughput of TCP flow from h1 to h2 by running:

```
mininet> h2 tail -f ./iperf-recv.txt
```

- You can quit viewing throughput by pressing CTRL-C.
- Think about the congestion window of the TCP stream again. How will it look for a long-established connection?
- You can also use ping to observe RTT while iperf is running

```
mininet> h1 ping -c 100 h2
```

#### The Impact on the Short Flow

- To see how our long-lived iperf flow affects our web page download, download the webpage again - while iperf is running. Observe how long it takes.

```
mininet> h2 wget h1
```

- Why does the web page take so much longer to download?

### 3.3 Measuring the real cwnd and buffer occupancy values.

We provided a script that lets you measure cwnd and buffer occupancy values (in terms of values for cwnd and buffer occupancy). We are going to re-run a couple of the experiments and plot the real values.

#### 1. Restart Mininet

- Stop and restart Mininet and the monitor script, then re-run the above experiment as follows.

```
mininet> exit
sudo ./run.sh
```

#### 2. Monitor TCP CWND and Buffer Occupancy in Mininet

- In your other terminal tab, type the following giving a name for your experiment.

```
./monitor.sh <EXP_NAME>
```

- In your first tab with the running mininet, start iperf again

```
mininet> h1 ./iperf.sh
```

- (wait for 70 seconds ...)

```
mininet> h2 wget h1
```

- Wait for the wget to complete, then stop the python monitor script followed by the instructions on the screen. The cwnd values are saved in <EXP\_NAME>\_tcpprobe.txt and the buffer occupancy in <EXP\_NAME>\_sw0-qlen.txt.

### 3. Plot CWND and Queue Occupancy

- Plot the TCP cwnd and queue occupancy from the output file

```
./plot_figures.sh <EXP_NAME>
```

- Adjust command line parameters to generate the figure you want.
- The script will also host a webserver on the machine and you can use the url the script provided to access to your figures.
- Note: the figures will sometimes interpolate where it is not appropriate. If you see a slow decrease of cwnd, then that probably is an artefact of such interpolation. Ignore those segments of the figures.
- If you are unable to see the cwnd, ensure you run wget after you started the monitor.sh script.

#### 3.4 Smaller Buffer: from 100 packets to 20 packets

Restart Mininet with small buffer

- Stop any running Mininet and start Mininet again, but this time we will make the buffers 20 packets long instead:

```
sudo ./run-minq.sh
```

- Repeat the earlier simple tests

```
mininet> h2 wget h1
mininet> h1 ping -c 10 h2
```

- Did the results change?
- Close mininet and start monitoring again

```
sudo ./monitor.sh <EXP_NAME>
```

- Restart mininet and re-run the experiment

```
mininet> h1 ./iperf.sh
mininet> h1 ping -c 30 h2
mininet> h2 wget h1
```

- What do you think the cwnd and queue occupancy will be like in this case?
- Plot the figure for cwnd and queue occupancy, this time using the script “./plot\_figures\_minq.sh”
- Then again, use the url to see your figures. Why does reducing the queue size reduce the download time for wget?

```
./plot_figures_minq.sh
```

## 4 UDP Experiments

- In this part, you write a simple UDP client/server system
- We are providing `client.py` and `server.py` scripts for your convenience.
  - We expect that `client.py` can be called with `-r` parameter, and we already provide the code for doing it in the attached client script.
- A helpful tutorial on a similar application: <https://pymotw.com/2/socket/udp.html>

### 4.1 Simple UDP client/server application using sockets

Your UDP server/client should follow these requirements:

- The server will receive incoming datagrams on port 5555
- The exact message format is up to you, it does not need to be efficient
- The client should send messages with a leading segment ID (e.g. part of the UDP payload)
- Your client should be configurable to send a fix amount of data per second (e.g. 1 Mbps), including UDP and IP header
- The server should verify if incoming traffic has missing datagrams, and display a warning if that is the case

### 4.2 Simple UDP connection

- start up mininet

```
sudo ./run.sh
```

- In the running mininet emulation, open a xterm window for h1, and for h2

```
mininet> xterm h1 h2
```

- Start your server application on h2, and the client on h1. The client should now be able to send data to the server.
- Set up your client to send approximately 1.5 Mbps of traffic to h2. Do you see dropped packets at the server?
- Set up your client to send a bit more, e.g. 1.6 Mbps. Again, do you see dropped packets?

## 5 What to Hand in

### 5.1 eDimension submission:

- Please provide a writeup (1-2 pages, PDF format with your name) that includes the following information:

- What is the normal time required to download the webpage on h1 from h2?
- What was your initial expectation for the congestion window size over time?
- After starting iperf on h1, did you observe something interesting in the ping RTT?
- After starting iperf on h1, why does the web page take so much longer to download?
- Please provide the figures for the first experiment (with qlen 100 and only one queue)
  - \* Please comment on what you can see in the figures
- Please provide the figures for the second experiment (with qlen 20 and only one queue)
  - \* Please comment on what you can see in the figures, and what is different (and why)
- Submit your server and client python code for the UDP part
  - The client should send approximately 1.5 Mbps
  - No packets should be dropped by default
  - If we increase the packet rate by 10%, we should see dropped packets
  - The server should report if incoming traffic has dropped datagrams (and, optionally, if a new sequence of data is started)
  - You are allowed to collaborate with another student, please specify both names and IDs in the client and server script files.