

Gede Ria Ghosalya
1001841

50.012 Networks – Lab 5

3.1 Warming Up

What is the local network chosen for the host?

```
mininet> h2 ifconfig
h2-eth0  Link encap:Ethernet  HWaddr 00:00:00:00:00:06
         inet addr:10.0.0.106 Bcast:10.0.0.255 Mask:255.255.255.0
         inet6 addr: fe80::200:ff:fe00:6/64 Scope:Link
         UP BROADCAST RUNNING MULTICAST  MTU:1500  Metric:1
         RX packets:114 errors:0 dropped:0 overruns:0 frame:0
         TX packets:14 errors:0 dropped:0 overruns:0 carrier:0
         collisions:0 txqueuelen:1000
         RX bytes:15423 (15.4 KB)  TX bytes:1848 (1.8 KB)

lo       Link encap:Local Loopback
         inet addr:127.0.0.1 Mask:255.0.0.0
         inet6 addr: ::1/128 Scope:Host
         UP LOOPBACK RUNNING  MTU:65536  Metric:1
         RX packets:0 errors:0 dropped:0 overruns:0 frame:0
         TX packets:0 errors:0 dropped:0 overruns:0 carrier:0
         collisions:0 txqueuelen:1
         RX bytes:0 (0.0 B)  TX bytes:0 (0.0 B)
```

Ans: 10.0.0.0/24

Are the two servers srv1 and srv2 in the same subnet?

```
mininet> srv1 route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.0.1 0.0.0.0 UG 0 0 0 srv1-eth0
10.0.0.0 0.0.0.0 255.0.0.0 U 0 0 0 srv1-eth0

mininet> srv2 route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.0.1 0.0.0.0 UG 0 0 0 srv2-eth0
10.0.0.0 0.0.0.0 255.0.0.0 U 0 0 0 srv2-eth0
```

Since both servers have the same IP after mask (subnet address), they are in the same subnet.

Test if you can tracepath from h1 to srv1.

```
mininet> h1 tracepath 10.0.0.10 -n
1?: [LOCALHOST] pmtu 1500
1: 10.0.0.10 1.257ms reached
1: 10.0.0.10 1.148ms reached
Resume: pmtu 1500 hops 1 back 1
```

We are unable to observe the switch; the only machine visible through tracepath is the server.

What is the gateway for all the devices?

```
mininet> h2 route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.0.111 0.0.0.0 UG 0 0 0 h2-eth0
10.0.0.0 0.0.0.0 255.255.255.0 U 0 0 0 h2-eth0
mininet> srv1 route -n
Kernel IP routing table
Destination Gateway Genmask Flags Metric Ref Use Iface
0.0.0.0 10.0.0.1 0.0.0.0 UG 0 0 0 srv1-eth0
10.0.0.0 0.0.0.0 255.0.0.0 U 0 0 0 srv1-eth0
```

Host gateway: 10.0.0.111

Server gateway: 10.0.0.1

Can you ping the server nils.net (8.8.8.2) from h1? If not, any idea what is going wrong?

```
mininet> h1 ping 8.8.8.2
PING 8.8.8.2 (8.8.8.2) 56(84) bytes of data.
From 10.0.0.105 icmp_seq=1 Destination Host Unreachable
From 10.0.0.105 icmp_seq=2 Destination Host Unreachable
From 10.0.0.105 icmp_seq=3 Destination Host Unreachable
```

H1 is unable to ping nils.net (8.8.8.2). No idea~

Is a DHCP server running in the local network? On which machine?

Source	Destination	Protocol
0.0.0.0	255.255.255.255	DHCP
10.0.0.10	10.0.0.105	DHCP
fa80::74c6:a1ff:fe01	ff02::fb	MDNS

```
mininet> h1 dhclient h1-eth0
RTNETLINK answers: File exists
mininet> h1 ifconfig
h1-eth0 Link encap:Ethernet HWaddr 00
inet addr:10.0.0.105 Bcast:10
inet6 addr: fe80::200:ff:fe00:
UP BROADCAST RUNNING MULTICAST
```

Yes, the DHCP is running on local network, which is on srv1.

4.1 Changing DHCP Configuration

Improve the configuration of the DHCP server.

```
1 interface=srv1-eth0
2 dhcp-range=srv1-eth0,10.0.0.100,10.0.0.200,255.255.255.0,12h
3 dhcp-option=3,10.0.0.111
4 dhcp-option=option:dns-server,0.0.0.0,8.8.8.8
5 dhcp-authoritative
```

On line 3, this configuration will have 10.0.0.111 as its router (option 3). However, we don't have a device with this address. Neither srv1 nor h1 can ping to this address.

```
mininet> srv1 ping 10.0.0.111
PING 10.0.0.111 (10.0.0.111) 56(84) bytes of data:
From 10.0.0.10 icmp_seq=1 Destination Host Unreachable
From 10.0.0.10 icmp_seq=2 Destination Host Unreachable
From 10.0.0.10 icmp_seq=3 Destination Host Unreachable
^C
--- 10.0.0.111 ping statistics ---
5 packets transmitted, 0 received, +3 errors, 100% packet loss, time 3999ms
pipe 3
mininet> h1 ping 10.0.0.111
PING 10.0.0.111 (10.0.0.111) 56(84) bytes of data:
From 10.0.0.105 icmp_seq=1 Destination Host Unreachable
From 10.0.0.105 icmp_seq=2 Destination Host Unreachable
From 10.0.0.105 icmp_seq=3 Destination Host Unreachable
^C
--- 10.0.0.111 ping statistics ---
4 packets transmitted, 0 received, +3 errors, 100% packet loss, time 2999ms
pipe 4
```

However, we do know that 10.0.0.1 exists as a router. So we can improve the configuration by changing the router.

```
1 interface=srv1-eth0
2 dhcp-range=srv1-eth0,10.0.0.100,10.0.0.200,255.255.255.0,12h
3 dhcp-option=3,10.0.0.1
4 dhcp-option=option:dns-server,0.0.0.0,8.8.8.8
5 dhcp-authoritative
```

Now h1 is able to ping 8.8.8.8

```
*** Starting CLI:
mininet> h1 ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data:
64 bytes from 8.8.8.8: icmp_seq=1 ttl=62 time=5.84 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=62 time=0.289 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=62 time=0.061 ms
64 bytes from 8.8.8.8: icmp_seq=4 ttl=62 time=0.059 ms
64 bytes from 8.8.8.8: icmp_seq=5 ttl=62 time=0.058 ms
```

4.2 DNS

On h1, try to ping nils.net. Can you reach it? Why?

```
mininet> h1 ping nils.net
PING nils.net (8.8.8.2) 56(84) bytes of data.
64 bytes from 8.8.8.2: icmp_seq=1 ttl=62 time=1.36 ms
64 bytes from 8.8.8.2: icmp_seq=2 ttl=62 time=0.284 ms
64 bytes from 8.8.8.2: icmp_seq=3 ttl=62 time=0.053 ms
64 bytes from 8.8.8.2: icmp_seq=4 ttl=62 time=0.065 ms
64 bytes from 8.8.8.2: icmp_seq=5 ttl=62 time=0.056 ms
```

Yes, h1 was able to ping nils.net, with the ip (8.8.8.2).

```
mininet> h1 dig nils.net
; <<>> DiG 9.10.3-P4-Ubuntu <<>> nils.net
;; global options: +cmd
;; Got answer:
;; ->>HEADER<<- opcode: QUERY, status: NOERROR, id: 46891
;; flags: qr aa rd ra ad; QUERY: 1, ANSWER: 1, AUTHORITY: 0, ADDITIONAL: 0

;; QUESTION SECTION:
;nils.net.                IN      A

;; ANSWER SECTION:
nils.net.                 0       IN      A      8.8.8.2

;; Query time: 2 msec
;; SERVER: 8.8.8.8#53(8.8.8.8)
;; WHEN: Fri Oct 13 11:19:21 SGT 2017
;; MSG SIZE rcvd: 42
```

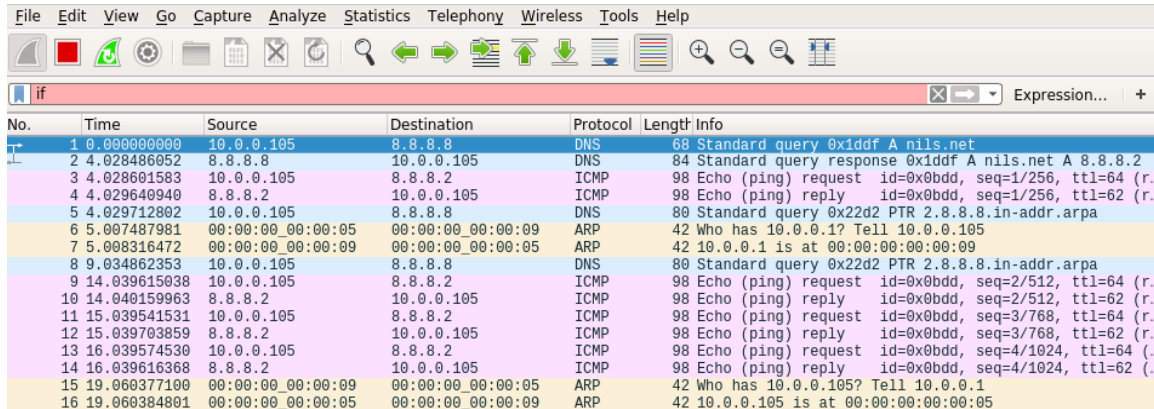
Upon using dig, it is discovered that the server 8.8.8.8 recognizes the address “nils.net” and return the IP address (8.8.8.2). As we know that h1 is able to reach 8.8.8.8, this is why h1 can connect to nils.net and recognize its IP address.

4.3 Observing NAT in action

In the provided setup, one node provides NAT for the hosts with private IP address. Which node is this?

The node providing NAT is intGW.

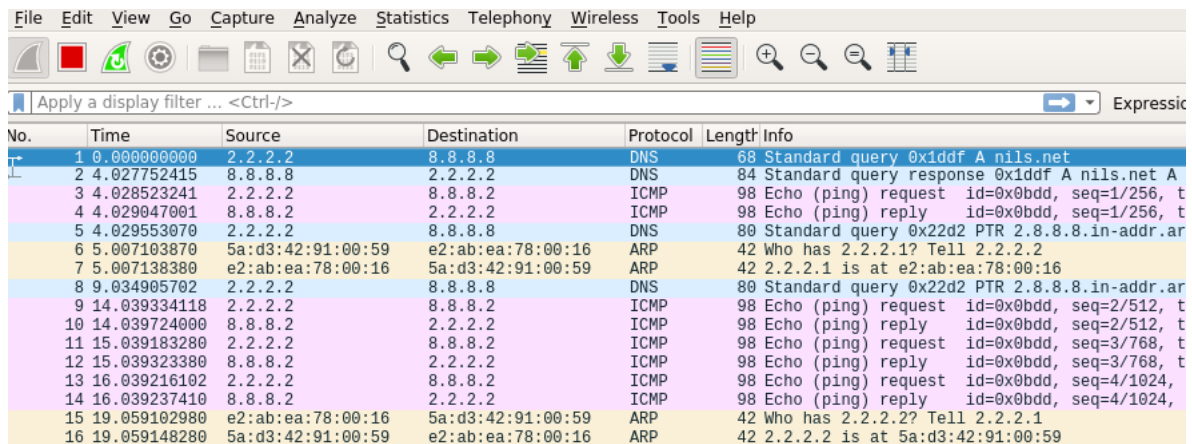
Below is 2 screen capture from wireshark, for h1 and intGW.



The screenshot shows a Wireshark capture from host h1 (10.0.0.105). The packet list contains 16 entries. The first packet is a DNS standard query from 10.0.0.105 to 8.8.8.8. The second packet is a DNS standard query response from 8.8.8.8 to 10.0.0.105. The third and fourth packets are ICMP echo (ping) request and reply between 10.0.0.105 and 8.8.8.2. The fifth and sixth packets are DNS standard query and response for PTR records between 10.0.0.105 and 8.8.8.8. The seventh and eighth packets are ARP requests and replies between 10.0.0.105 and 8.8.8.2. The ninth and tenth packets are ICMP echo (ping) request and reply between 10.0.0.105 and 8.8.8.2. The eleventh and twelfth packets are ICMP echo (ping) request and reply between 10.0.0.105 and 8.8.8.2. The thirteenth and fourteenth packets are ICMP echo (ping) request and reply between 10.0.0.105 and 8.8.8.2. The fifteenth and sixteenth packets are ARP requests and replies between 10.0.0.105 and 8.8.8.2.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	10.0.0.105	8.8.8.8	DNS	68	Standard query 0x1dddf A nils.net
2	4.028486052	8.8.8.8	10.0.0.105	DNS	84	Standard query response 0x1dddf A nils.net A 8.8.8.2
3	4.028601583	10.0.0.105	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=1/256, ttl=64 (r...
4	4.029640940	8.8.8.2	10.0.0.105	ICMP	98	Echo (ping) reply id=0x0bdd, seq=1/256, ttl=62 (r...
5	4.029712882	10.0.0.105	8.8.8.8	DNS	80	Standard query 0x22d2 PTR 2.8.8.8.in-addr.arpa
6	5.007487981	00:00:00_00:00:05	00:00:00_00:00:09	ARP	42	Who has 10.0.0.1? Tell 10.0.0.105
7	5.008316472	00:00:00_00:00:09	00:00:00_00:00:05	ARP	42	10.0.0.1 is at 00:00:00:00:00:09
8	9.034862353	10.0.0.105	8.8.8.8	DNS	80	Standard query 0x22d2 PTR 2.8.8.8.in-addr.arpa
9	14.039615038	10.0.0.105	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=2/512, ttl=64 (r...
10	14.040159963	8.8.8.2	10.0.0.105	ICMP	98	Echo (ping) reply id=0x0bdd, seq=2/512, ttl=62 (r...
11	15.039541531	10.0.0.105	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=3/768, ttl=64 (r...
12	15.039703859	8.8.8.2	10.0.0.105	ICMP	98	Echo (ping) reply id=0x0bdd, seq=3/768, ttl=62 (r...
13	16.039574530	10.0.0.105	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=4/1024, ttl=64 (r...
14	16.039616368	8.8.8.2	10.0.0.105	ICMP	98	Echo (ping) reply id=0x0bdd, seq=4/1024, ttl=62 (r...
15	19.060377100	00:00:00_00:00:09	00:00:00_00:00:05	ARP	42	Who has 10.0.0.105? Tell 10.0.0.1
16	19.060384801	00:00:00_00:00:05	00:00:00_00:00:09	ARP	42	10.0.0.105 is at 00:00:00:00:00:05

Figure 4.3.1 wireshark screen capture from h1



The screenshot shows a Wireshark capture from the intGW node. The packet list contains 16 entries. The first packet is a DNS standard query from 2.2.2.2 to 8.8.8.8. The second packet is a DNS standard query response from 8.8.8.8 to 2.2.2.2. The third and fourth packets are ICMP echo (ping) request and reply between 2.2.2.2 and 8.8.8.2. The fifth and sixth packets are DNS standard query and response for PTR records between 2.2.2.2 and 8.8.8.8. The seventh and eighth packets are ARP requests and replies between 2.2.2.2 and 8.8.8.2. The ninth and tenth packets are ICMP echo (ping) request and reply between 2.2.2.2 and 8.8.8.2. The eleventh and twelfth packets are ICMP echo (ping) request and reply between 2.2.2.2 and 8.8.8.2. The thirteenth and fourteenth packets are ICMP echo (ping) request and reply between 2.2.2.2 and 8.8.8.2. The fifteenth and sixteenth packets are ARP requests and replies between 2.2.2.2 and 8.8.8.2.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000000	2.2.2.2	8.8.8.8	DNS	68	Standard query 0x1dddf A nils.net
2	4.027752415	8.8.8.8	2.2.2.2	DNS	84	Standard query response 0x1dddf A nils.net A
3	4.028523241	2.2.2.2	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=1/256, t
4	4.029047001	8.8.8.2	2.2.2.2	ICMP	98	Echo (ping) reply id=0x0bdd, seq=1/256, t
5	4.029553070	2.2.2.2	8.8.8.8	DNS	80	Standard query 0x22d2 PTR 2.8.8.8.in-addr.ar
6	5.007103870	5a:d3:42:91:00:59	e2:ab:ea:78:00:16	ARP	42	Who has 2.2.2.1? Tell 2.2.2.2
7	5.007138380	e2:ab:ea:78:00:16	5a:d3:42:91:00:59	ARP	42	2.2.2.1 is at e2:ab:ea:78:00:16
8	9.034905702	2.2.2.2	8.8.8.8	DNS	80	Standard query 0x22d2 PTR 2.8.8.8.in-addr.ar
9	14.039334118	2.2.2.2	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=2/512, t
10	14.039724000	8.8.8.2	2.2.2.2	ICMP	98	Echo (ping) reply id=0x0bdd, seq=2/512, t
11	15.039183280	2.2.2.2	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=3/768, t
12	15.039323380	8.8.8.2	2.2.2.2	ICMP	98	Echo (ping) reply id=0x0bdd, seq=3/768, t
13	16.039216102	2.2.2.2	8.8.8.2	ICMP	98	Echo (ping) request id=0x0bdd, seq=4/1024, t
14	16.039237410	8.8.8.2	2.2.2.2	ICMP	98	Echo (ping) reply id=0x0bdd, seq=4/1024, t
15	19.059102980	e2:ab:ea:78:00:16	5a:d3:42:91:00:59	ARP	42	Who has 2.2.2.2? Tell 2.2.2.1
16	19.059148280	5a:d3:42:91:00:59	e2:ab:ea:78:00:16	ARP	42	2.2.2.2 is at 5a:d3:42:91:00:59

Figure 4.3.2 wireshark screen capture from intGW

It can be seen that the request sent from h1 (10.0.0.105) to nils.net is reflected as request from intGW (2.2.2.2) to nils.net.

Opening net.py, the *enableNAT* function is written as follows.

```
def enableNAT(net,hostn):
    # this assumes that internal interface is eth0, external interface is eth1, and the network is 10.0.0.0/24
    host = net.getNodeByName(hostn)
    host.cmd("iptables -A FORWARD -o %s-eth1 -i %s-eth0 -s 10.0.0.0/24 -m conntrack --ctstate NEW -j ACCEPT"%(hostn,hostn))
    host.cmd("iptables -A FORWARD -m conntrack --ctstate ESTABLISHED,RELATED -j ACCEPT")
    host.cmd("iptables -t nat -F POSTROUTING")
    host.cmd("iptables -t nat -A POSTROUTING -o %s-eth1 -j MASQUERADE"%hostn)
```

From here we can see that the address range for NAT translation is 10.0.0.0/24 (or 10.0.0.0 10.0.0.255)

4.4 Simple Firewalling

Initially, srv2 (10.0.0.11) is able to ping both intGW (10.0.0.1) and outside network (8.8.8.8) as shown below.

```
root@desktop:~/gede/networks_lab/lab5# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
4 bytes from 8.8.8.8: icmp_seq=1 ttl=62 time=3.02 ms
4 bytes from 8.8.8.8: icmp_seq=2 ttl=62 time=0.251 ms
^C
-- 8.8.8.8 ping statistics ---
  packets transmitted, 2 received, 0% packet loss, time 1001ms
  tt min/avg/max/mdev = 0.251/1.639/3.027/1.388 ms
root@desktop:~/gede/networks_lab/lab5# ping 10.0.0.1
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
4 bytes from 10.0.0.1: icmp_seq=1 ttl=64 time=1.33 ms
4 bytes from 10.0.0.1: icmp_seq=2 ttl=64 time=0.244 ms
^C
-- 10.0.0.1 ping statistics ---
  packets transmitted, 2 received, 0% packet loss, time 1001ms
  tt min/avg/max/mdev = 0.244/0.789/1.334/0.545 ms
```

However, we are able to block srv2 by adding REJECT rule to iptables, as follows.

```
root@desktop:~/gede/networks_lab/lab5# iptables -I INPUT -s 10.0.0.11 -j REJECT
root@desktop:~/gede/networks_lab/lab5# iptables -I OUTPUT -s 10.0.0.11 -j REJECT
root@desktop:~/gede/networks_lab/lab5# iptables -I FORWARD -s 10.0.0.11 -j REJECT
root@desktop:~/gede/networks_lab/lab5#
```

Note that FORWARD rule should also be added to prevent forward connection from srv2 to outside network, such as 8.8.8.8 node. After the above rule is added, we can see that srv2 is no longer able to ping neither intGW nor 8.8.8.8 node.

```
root@desktop:~/gede/networks_lab/lab5# ping 10.0.0.1
PING 10.0.0.1 (10.0.0.1) 56(84) bytes of data.
From 10.0.0.1 icmp_seq=1 Destination Port Unreachable
From 10.0.0.1 icmp_seq=2 Destination Port Unreachable
^C
--- 10.0.0.1 ping statistics ---
 2 packets transmitted, 0 received, +2 errors, 100% packet loss, time 999ms

root@desktop:~/gede/networks_lab/lab5# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
From 10.0.0.1 icmp_seq=1 Destination Port Unreachable
From 10.0.0.1 icmp_seq=2 Destination Port Unreachable
^C
--- 8.8.8.8 ping statistics ---
 2 packets transmitted, 0 received, +2 errors, 100% packet loss, time 999ms
```